

ty6dmp17x

April 8, 2024

```
[1]: # Aim: To perform and find the accuracy of Logistic Regression
```

```
[2]: # Name: Mandar K Satpute  
# Sub: Big Data Analytics (ET-Lab 2)  
# Section : B  
# Roll no: 54
```

```
[3]: import pandas as pd
```

```
[4]: from sklearn.model_selection import train_test_split
```

```
[5]: from sklearn.linear_model import LogisticRegression
```

```
[6]: from sklearn.metrics import accuracy_score
```

```
[7]: df=pd.read_csv('C:\\Users\\hp\\Desktop\\CHD_preprocessed.csv')
```

```
[8]: df.head()
```

```
[8]:
```

	male	age	education	currentSmoker	cigsPerDay	BPMeds	prevalentStroke	\
0	1	39	1	0	0.0	0.0	0	
1	0	46	0	0	0.0	0.0	0	
2	1	48	0	1	20.0	0.0	0	
3	0	61	1	1	30.0	0.0	0	
4	0	46	1	1	23.0	0.0	0	

  

	prevalentHyp	diabetes	totChol	sysBP	diaBP	BMI	heartRate	glucose	\
0	0	0	195.0	106.0	70.0	26.97	80.0	77.0	
1	0	0	250.0	121.0	81.0	28.73	95.0	76.0	
2	0	0	245.0	127.5	80.0	25.34	75.0	70.0	
3	1	0	225.0	150.0	95.0	28.58	65.0	103.0	
4	0	0	285.0	130.0	84.0	23.10	85.0	85.0	

  

	TenYearCHD
0	0
1	0
2	0

```
3          1
4          0
```

```
[9]: df.tail()
```

```
[9]:      male  age  education  currentSmoker  cigsPerDay  BPMeds  \
4128     1   50         0           1         1.0     0.0
4129     1   51         1           1        43.0     0.0
4130     0   48         0           1        20.0     0.0
4131     0   44         0           1        15.0     0.0
4132     0   52         0           0         0.0     0.0

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
4128                0             1         0    313.0  179.0   92.0  25.97
4129                0             0         0    207.0  126.5   80.0  19.71
4130                0             0         0    248.0  131.0   72.0  22.00
4131                0             0         0    210.0  126.5   87.0  19.16
4132                0             0         0    269.0  133.5   83.0  21.47

      heartRate  glucose  TenYearCHD
4128       66.0     86.0           1
4129       65.0     68.0           0
4130       84.0     86.0           0
4131       86.0     82.0           0
4132       80.0    107.0           0
```

```
[10]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4133 entries, 0 to 4132
Data columns (total 16 columns):
#   Column                Non-Null Count  Dtype
---  -
0   male                  4133 non-null  int64
1   age                   4133 non-null  int64
2   education              4133 non-null  int64
3   currentSmoker          4133 non-null  int64
4   cigsPerDay              4133 non-null  float64
5   BPMeds                 4133 non-null  float64
6   prevalentStroke         4133 non-null  int64
7   prevalentHyp            4133 non-null  int64
8   diabetes               4133 non-null  int64
9   totChol                4133 non-null  float64
10  sysBP                 4133 non-null  float64
11  diaBP                 4133 non-null  float64
12  BMI                   4133 non-null  float64
13  heartRate              4133 non-null  float64
```

```

14 glucose          4133 non-null   float64
15 TenYearCHD       4133 non-null   int64
dtypes: float64(8), int64(8)
memory usage: 516.8 KB

```

```
[11]: df.describe()
```

```

[11]:
count      male      age  education  currentSmoker  cigsPerDay \
mean      0.427293  49.557222   0.280668      0.494798      9.101621
std       0.494745   8.561628   0.449380      0.500033     11.918440
min       0.000000  32.000000   0.000000      0.000000      0.000000
25%       0.000000  42.000000   0.000000      0.000000      0.000000
50%       0.000000  49.000000   0.000000      0.000000      0.000000
75%       1.000000  56.000000   1.000000      1.000000     20.000000
max       1.000000  70.000000   1.000000      1.000000     70.000000

count      BPMeds  prevalentStroke  prevalentHyp      diabetes      totChol \
mean      0.034358      0.006049      0.311154      0.025647     236.664408
std       0.182168      0.077548      0.463022      0.158100     43.909188
min       0.000000      0.000000      0.000000      0.000000     107.000000
25%       0.000000      0.000000      0.000000      0.000000     206.000000
50%       0.000000      0.000000      0.000000      0.000000     234.000000
75%       0.000000      0.000000      1.000000      0.000000     262.000000
max       1.000000      1.000000      1.000000      1.000000     600.000000

count      sysBP      diaBP      BMI      heartRate      glucose \
mean     132.367046     82.872248     25.778571     75.925236     81.946528
std       22.080332     11.952654      4.074360     12.049188     22.860954
min       83.500000     48.000000     15.540000     44.000000     40.000000
25%      117.000000     75.000000     23.060000     68.000000     72.000000
50%      128.000000     82.000000     25.380000     75.000000     80.000000
75%      144.000000     89.500000     27.990000     83.000000     85.000000
max      295.000000    142.500000     56.800000    143.000000    394.000000

count      TenYearCHD
mean      0.151948
std       0.359014
min       0.000000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.000000

```

```
[12]: df.isnull()
```

```
[12]:      male    age  education  currentSmoker  cigsPerDay  BPMeds  \
0      False  False      False      False      False    False
1      False  False      False      False      False    False
2      False  False      False      False      False    False
3      False  False      False      False      False    False
4      False  False      False      False      False    False
...
4128   False  False      False      False      False    False
4129   False  False      False      False      False    False
4130   False  False      False      False      False    False
4131   False  False      False      False      False    False
4132   False  False      False      False      False    False

      prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0              False      False    False    False  False  False  False
1              False      False    False    False  False  False  False
2              False      False    False    False  False  False  False
3              False      False    False    False  False  False  False
4              False      False    False    False  False  False  False
...
4128           False      False    False    False  False  False  False
4129           False      False    False    False  False  False  False
4130           False      False    False    False  False  False  False
4131           False      False    False    False  False  False  False
4132           False      False    False    False  False  False  False

      heartRate  glucose  TenYearCHD
0           False    False      False
1           False    False      False
2           False    False      False
3           False    False      False
4           False    False      False
...
4128           False    False      False
4129           False    False      False
4130           False    False      False
4131           False    False      False
4132           False    False      False
```

```
[4133 rows x 16 columns]
```

```
[13]: df.isna().sum()
```

```
[13]: male          0
      age          0
```

```

education      0
currentSmoker  0
cigsPerDay     0
BPMeds         0
prevalentStroke 0
prevalentHyp   0
diabetes       0
totChol        0
sysBP          0
diaBP          0
BMI            0
heartRate      0
glucose        0
TenYearCHD     0
dtype: int64

```

```

[14]: x=df.drop("TenYearCHD",axis=1)
      y=df['TenYearCHD']

```

```

[15]: x

```

```

[15]:
   male  age  education  currentSmoker  cigsPerDay  BPMeds  \
0      1   39         1             0         0.0     0.0
1      0   46         0             0         0.0     0.0
2      1   48         0             1        20.0     0.0
3      0   61         1             1        30.0     0.0
4      0   46         1             1        23.0     0.0
...    ...   ...     ...             ...         ...     ...
4128    1   50         0             1         1.0     0.0
4129    1   51         1             1        43.0     0.0
4130    0   48         0             1        20.0     0.0
4131    0   44         0             1        15.0     0.0
4132    0   52         0             0         0.0     0.0

   prevalentStroke  prevalentHyp  diabetes  totChol  sysBP  diaBP  BMI  \
0                0             0         0    195.0   106.0   70.0  26.97
1                0             0         0    250.0   121.0   81.0  28.73
2                0             0         0    245.0   127.5   80.0  25.34
3                0             1         0    225.0   150.0   95.0  28.58
4                0             0         0    285.0   130.0   84.0  23.10
...              ...           ...         ...         ...         ...
4128              0             1         0    313.0   179.0   92.0  25.97
4129              0             0         0    207.0   126.5   80.0  19.71
4130              0             0         0    248.0   131.0   72.0  22.00
4131              0             0         0    210.0   126.5   87.0  19.16
4132              0             0         0    269.0   133.5   83.0  21.47

```

	heartRate	glucose
0	80.0	77.0
1	95.0	76.0
2	75.0	70.0
3	65.0	103.0
4	85.0	85.0
...	...	...
4128	66.0	86.0
4129	65.0	68.0
4130	84.0	86.0
4131	86.0	82.0
4132	80.0	107.0

[4133 rows x 15 columns]

[16]: y

```
[16]: 0      0
      1      0
      2      0
      3      1
      4      0
      ..
      4128    1
      4129    0
      4130    0
      4131    0
      4132    0
      Name: TenYearCHD, Length: 4133, dtype: int64
```

[17]: x\_train,x\_test,y\_train,y\_test= train\_test\_split(x,y,test\_size=0.  
↪2,random\_state=42)

[18]: model = LogisticRegression(max\_iter=1600)

[19]: model.fit(x\_train,y\_train)  
model.score(x\_train, y\_train)

[19]: 0.8623714458560193